The guide to defining a custom Xillybus IP core

1

Introduction

Xillybus is a multi-purpose platform for a variety of applications. As such, the delivered IP core is easily configured to meet specific requirements in terms of the number of streams, their direction, attributes related to their performance and consumption of resources.

Custom IPs are generated for evaluation and licensing alike. Xillybus' "try it first" policy encourages potential licensees to request a tailored version of the Xillybus IP core for evaluation under real-life conditions.

To simplify the definition and reception of custom IP cores, an online tool for instant generation is available at http://xillybus.com/custom-ip-factory

This tool consists of a simple web interface, which allows the user to define the device files requested and their configuration. Once the definition is complete and submitted, an automatic process generates the bundle for inclusion in the FPGA project. The custom IP core is ready for download after a short while, typically a few minutes.

The web interface may be used without reading this guide, but it's recommended to first familiarize yourself with Xillybus by running the demo bundle. Users who wish to get a better understanding and control of the device files' attributes, and set them manually by disabling the "autoset" option, will find some background information in this guide.

For users who have not yet familiarized themselves with the demo bundle, the following documents are recommended for prior reading:

- Getting started with the FPGA demo bundle for Xilinx

- Getting started with the FPGA demo bundle for Intel FPGA

- Getting started with Xillinux for Zynq-7000

- Getting started with Xillybus on a Linux host

- Getting started with Xillybus on a Windows host

- The guide to Xillybus Block Design Flow for non-HDL users

Even when the need for a custom IP core is clear, it's best to start off with the demo bundle's IP core and project, as it shows how the IP core is integrated with the application logic, and how the entire project should be set up.

All information about the IP core's custom configuration is stored within the IP core itself in the FPGA. The host driver retrieves this information at its initialization, so there is no need to change anything on the host when switching from one Xillybus IP configuration to another.

A common mistake is trying to minimize the number of streams configured, for the sake of saving logic resources. Sections 3 and 4.3 show how a Xillybus IP core scales, and explains why it makes sense allocating streams generously.

2

Defining custom IP cores

2.1   Overview

The web tool provides a wizard-like interface for defining a custom IP core from scratch, or using the configuration of the demo bundle's core as a starting point.

There are a few points worth emphasizing when using this tool:

- For most purposes, it's recommended to keep the "Autoset internals" option selected when defining device files. Manual setting of the data flow control and buffer attributes is likely to have a negative effect on performance.

- It's important to set the target FPGA device family correctly, since the IP core is delivered as an architecture-dependent netlist binary.

- It's also important to set each device file's "use" attribute to the description best matching the intended purpose.

- For XillyUSB IP cores, the "Expected bandwidth" attribute should be set accurately to the maximally requested bandwidth by the stream, as the data rate is limited to that value. For other variants (PCIe and AXI), this attribute only affects performance tuning. Realistic ballpark figures should be applied, rather than attempting to "push the tool" by exaggerating the requirements. Such exaggeration may result in a performance hit on other streams that really need certain limited resources.

The rest of this section discusses some of the device files' attributes.

2.2   The device file's name

Each stream is designated a name, which appears in the host environment for its unique identifications.

The names always take the form xillybus_*, e.g. xillybus_mystream. Or for XillyUSB IP cores, xillyusb_NN_*, where NN is an index – typically two zeros when only one XillyUSB device is connected to the host.

On a Linux system, the stream is opened as a the plain file, e.g. /dev/xillybus_mystream. In Windows, the same stream appears as \\.\xillybus_mystream.

A device file can represent two streams in opposite directions, which is just two streams happening to share a device file name. These two streams can be opened separately in either direction, or opened for read-write. This feature should be avoided in general to prevent confusion, but is useful when the device file is passed to software expecting a bidirectional pipe.

2.3   Data width

The data width corresponds to the word fetched from or written to the FIFOs in the FPGA. The allowed choices are 32, 16 or 8 bits. Wider data widths are allowed by XillyUSB IP cores, as well as revision B/XL/XXL IP cores (the latter is discussed in section 4).

Streams requiring high bandwidth performance on Revision A (baseline) IP cores for PCIe and all IP cores for AXI, must be set to use 32 bits data width. There's a significant performance degradation for 16 and 8-bit data width, leading to inefficient use of the underlying transport (e.g. the PCIe bus transport).

The reason is that the words are transported through the Xillybus internal data paths at the bus clock rate. As a result, transporting an 8-bit word takes the same time slot as a 32-bit word, making it effectively four times slower.

This also impacts other streams competing for the underlying transport at a given time, since the data paths become occupied with slower data elements.

Later IP core revisions, as well as XillyUSB, have a different internal data path structure, which is why this doesn't apply to them.

Regardless, it's good practice to perform I/O operations in the host application with data width granularity, e.g. receive and send buffers with sizes that are a multiple of 4 if the data width is 32 bits.

A poor choice of data width may lead to undesired behavior. For example, if an host-to-FPGA link is 32 bits wide, writing 3 bytes of data at the host will make the driver wait, possibly forever, for the fourth byte before sending anything to the FPGA.

2.4 Synchronous or asynchronous stream

This attribute is set automatically when the "Autoset internals" option is selected, based upon the selection of the "use" setting.

In most cases, asynchronous streams are adequate for continuous data streams, and synchronous streams are adequate for control data.

For synchronous streams, all I/O (including the data flow in the FPGA) takes place only between the invocation and return of the system read() or write() call. This gives full control on what happens when, but leaves the stream unused while the CPU is doing other things. It's recommended to read the elaboration on this subject in section 2 of one of these two documents:

- Xillybus host application programming guide for Linux

- Xillybus host application programming guide for Windows

Synchronous streams make the programming more intuitive, but have a negative impact on bandwidth utilization. Asynchronous streams make it possible to maintain a continuous data flow, despite the operating system depriving user space processes from CPU for certain periods of time.

To summarize this subject, these are the guiding questions for either direction:

- For downstreams: Is it OK that a write() operation returns before the data has reached the FPGA?

- For upstreams: Is it OK that the Xillybus core begins fetching data from the user logic in the FPGA before a read() operation in the host requests it?

If the answer to the respective question is "no", a synchronous stream is needed. Otherwise, the asynchronous option is usually preferred, along with the understanding that the data flow is slightly less intuitive.

2.5 Buffering time

Xillybus maintains an illusion of a continuous stream of data between the FPGA and the host. The existence of driver's buffers is transparent to the user application logic

in the FPGA as well as the application on the host. They are of interest only to control the efficiency of the data flow and its ability to remain continuous, in particular at high data rates.

Data capture and playback applications require a continuous flow of data at the FPGA, or data is lost. To maintain this flow, a user-space application needs to make read() or write() calls frequently enough to prevent the driver's buffers from overflowing or underflowing (respectively).

Common operating systems, such as Linux and Windows, may deprive any user-space application from the CPU for theoretically arbitrary periods of time. The FPGA keeps filling or draining the driver's buffers during these periods of time, causing the data flow to stall, unless the buffers are large enough.

When setting a Xillybus stream for a use requiring continuity, and choosing "Autoset internals", a "Buffering" selection box appears in the web interface. The time periods chosen should reflect the expected maximal CPU deprivation period. Choosing "Maximum" tells the buffer allocation algorithm to attempt allocating as much RAM as possible, with just some consideration for the other streams.

Given a desired buffering time t and an expected bandwidth W, the ideal total RAM allocated for the driver's buffers of the stream, M, is

$$M = t \times W$$

The actual buffer sizes are however always a power of 2. It may also turn out impossible to allocate enough memory to meet the desired buffering time.

It is therefore important to look up the allocated buffer size in the core's readme file, and verify that it's acceptable to work with. Setting the buffer size manually (i.e. turning off "Autoset internals") may be necessary to force a distribution of buffer RAM more suitable for the target application.

2.6   Host driver's buffer size

It's recommended to let the tools set up the buffers' parameters automatically by enabling the "Autoset internals" in the web interface. See section 2.5 above.

The issue of driver's buffers is less significant for synchronous streams. For such, the rule of thumb is that the total buffer space allocated for a stream should be in the order of magnitude of the chunks of data for which data is transported. There is rarely any point in wasting kernel RAM space by making them larger than so.

For asynchronous streams, the buffers' parameters have a significant impact which is discussed in the section named "Continuous high rate I/O" in these two documents:

- Xillybus host application programming guide for Linux

- Xillybus host application programming guide for Windows

For the sake of continuity, the more buffer space the better, as the total buffer space keeps the streaming of data continuous even when the CPU is deprived from the application. Making a correct decision involves other factors, which are detailed in the programming guides referenced above.

A possible drawback of excessively large buffer space is a buffering delay, which is a result of the ability to store a large amount of data, so it arrives at the other end after a larger amount of time. This can be controlled by the technique mentioned in Xillybus FPGA designer's guide, in the section named "Monitoring the amount of buffered data".

For XillyUSB IP cores, there is one buffer for each stream, which functions as a large FIFO that is managed by the driver. Other IP cores (PCIe and AXI) maintain several DMA buffers for each stream, so both their size and count are defined. The effective buffer space is hence the size of each DMA buffer multiplied by their number.

The size of each underlying DMA buffer has a significance of its own in streams from the host to FPGA: The data is sent to the FPGA when these buffers are full (unless the buffer is flushed manually or by virtue of a timeout in the driver). The size of each DMA buffer has therefore an impact of the typical latency of flowing data.

2.7   DMA acceleration

Host to FPGA streams may sometimes require acceleration of the DMA data transfers, when the underlying transport is PCIe.

The PCIe bus protocol states that the FPGA should issue requests for data from the host and wait for the data to arrive. An inherent delay occurs as the request travels along the bus path, is queued and handled by the host, and the data travels back. This turnaround time gap causes a certain degradation in the bus efficiency, sometimes reducing the bandwidth of a single stream to as low as 40%.

To work around this issue, multiple data requests are sent, so that the host always has a request in its queue during continuous transmissions. Since data from different requests may arrive unordered, it must be stored in RAM buffers on the FPGA to present an ordered stream of data to the application logic.

Each buffer in the FPGA is used to store a segment of requested data. The current possible settings for DMA accelerations are

- None. No data is stored on the FPGA. Each request for data is sent only when all data has arrived from the previous one.

- 4 segments of 512 bytes each. 2048 bytes of buffer space is allocated on the FPGA. Up to four data requests can be "in flight" at any given moment.

- 8 segments of 512 bytes each. 4096 bytes of buffer space is allocated on the FPGA. Up to eight data requests can be "in flight" at any given moment.

- Revision B and later IP cores have an option of 16 segments of 512 bytes each as well.

The turnaround time between a request and the data arrival depends on the host's hardware. The actual bandwidth performance may therefore vary.

When using the IP Core Factory, the automatic allocation of acceleration resources is based upon measured results on typical PC hardware, and may need manual refinements in rare cases.

3

Scalability and logic resource consumption

3.1 General

Xillybus was designed with scalability in mind. While it makes perfect sense to configure a custom IP core for as little as a single stream, scaling up to a large number of streams has a relatively small impact on the amount of logic consumed by the Xillybus core.

In order to measure the logic resource consumption, successive builds of the Revision A (baseline) Xillybus IP core were made with an increasing number of streams. In all tests, the number of streams from the FPGA to the host were the same as in the other direction. The number of streams tested for ranged from 2 (one in each direction) to 64 (32 in each direction).

This section outlines the logic consumption of the IP core itself on three FPGA devices by Xilinx, as reported by their tools. The FPGAs targeted in the graphs that follow are quite outdated, however similar results are achieved on Intel devices and other Xilinx device families.

For a similar analysis of revision B, XL and XXL IP cores, see section 4.

XillyUSB is not covered in any of these analyses.

3.2 Block RAMs

The number of block RAMs used by the Xillybus core varies between zero to a few of them (3 block RAMs for 64 streams). There are no per-stream buffers inside the Xillybus core. Rather, the Xillybus core relies on the FIFOs connected to it to collect the data. Internally, the core has a single pool of memory used by all streams, storing

data for immediate transmission.

As the number of streams grow, block RAMs are chosen for efficient storage of host DMA addresses.

Further block RAMs are allocated for DMA acceleration of host to FPGA streams as required and detailed in the core's readme file.

3.3   Logic fabric resources

The graphs below show the consumption of LUTs and registers (flip-flops) as the number of streams go from 2 to 64. Each dot in those graphs is the de-facto use resulting from the synthesis report. What is evident from these graphs is the nearly linear growth in logic consumption. Regardless of the FPGA architecture, each stream adds about 110 LUTs and 82 registers on the average.

The number of actual slices consumed on the FPGA depends on how well the logic elements are packed into them. Each slice in Spartan-6 or Virtex-6 FPGAs can contain up to 8 LUTs and 8 registers. Accordingly, a very optimistic approach would be to assume that the registers are packed perfectly, so each stream adds only 110/8 =14 slices to the design. On the other hand, packing with half that efficiency is something achievable with no considerable effort. So the expected cost in slices for a stream can be estimated in the range of 14-28 slices.

It's important to note that when the FPGA device isn't nearly full, the Xilinx implementation tools don't bother to pack logic into slices efficiently, so the increase in slice count can be significantly steeper. This is merely a waste of resources there's plenty of anyhow.

The chosen setting for the benchmark test was 50% upstreams and 50% downstreams. Real-life IP cores usually have an emphasis on either directions, but the results below give an idea of what to expect.

The graphs follow. The slope may appear steep, but note that the stream count goes from a minimal design (2 streams) to a rather heavy one (64 streams).

The bottom line is that it makes sense to allocate extra streams in the design, even for the most trivial tasks, since their slice count contribution is fairly low.

Register count (synthesis)

LUT count (synthesis)

4

Revision B, XL and XXL IP cores

4.1 General

Up to this point, this document has related to the revision A (baseline) IP core, which is available since 2010. Revision B and XL cores were introduced in 2015, adapting to de-facto needs of Xillybus' user base. These cores will gradually replace Revision A cores.

Revision XXL cores were introduced in 2019, mostly because it was the only way to handle Intel's Stratix 10 FPGAs. Even though they supply double the bandwidth compared with XL, this wasn't the reason behind their introduction.

The new revisions (B, XL and XXL) offer a superset of features compared with revision A, but are functionally equivalent when defined with the same attributes (with some possible performance improvements).

The most notable differences are:

- Increased data bandwidth: Revision B cores' aggregate bandwidth limit is approximately double of its revision A counterpart; revision XL cores supply about four times as much bandwidth, compared with revision A. XXL about eight times.

- User interface data widths of 64, 128, and 256 bits are allowed in addition to the already existing options of 8, 16 and 32 bits. These widths are allowed regardless of the width of the data paths between Xillybus' IP core and the PCIe block in use.

- Faster logic design (easier to meet timing constraints), about 1 ns better on the slowest path.

- The logic resources consumption is lower in most common cases (see section

4.3).

- The PCIe bandwidth utilization remains efficient regardless of the user interface data width on revisions B, XL and XXL. This is contrary to the lower efficiency regarding widths 8 and 16 bits on revision A.

- On Xilinx platforms, Revision B, XL and XXL are available only for the Vivado toolchain.

4.2 Matching demo bundles

Revision B IP cores are drop-in replacements for revision A cores. Hence, the baseline demo bundle for the desired FPGA target should be used as a starting point. As this demo bundle arrives with a revision A core, those who desire a revision B core should configure and download it from the IP Core Factory.

Revision XL and XXL IP cores, on the other hand, require a dedicated demo bundle to work against. It may be needed to request such demo bundles by email.

4.3 Logic resource consumption

Revision B/XL/XXL IP cores are optimized for speed and a slightly lower logic consumption, at the cost of a slightly steeper logic consumption as the number of streams increases.

In order to quantify the use of logic resources, cores with an increasing number of streams were generated, targeting Kintex-7. The cores were synthesized, and the logic elements counted. As in section 3.3, the benchmark test was 50% upstreams and 50% downstreams.

The following three charts compare the logic consumption of Revision A, B and XL IP cores with equal settings. All tested streams were 32 bits wide.

Counting registers and LUTs, revision B cores outperform revision A cores when the stream count is low, but lose this advantage as the streams mount up.

Revision XL and XXL cores consume more logic than both other revisions in all scenarios.

The chart for block RAMs shows that both revision B and XL consume double as many block RAMs, compared with revision A.

The suggested conclusion is that revision B should almost always be preferred over revision A. Even when the logic count comparison says the opposite, does the im-

proved timing of revision B outweight the difference of logic use, in particular as it's negligible compared with the FPGA's capacity in most practical scenarios.

Revision XL and XXL, on the other hand, should be chosen only for applications that require their respective bandwidth capacity, as they are heavier in terms of logic consumption as well as timing.

LUT count (synthesis)

Block RAM count (synthesis)

4.4 Tuning for optimal Host-to-FPGA bandwidth

The higher data rates of revision B, XL and XXL cores makes the performance increasingly sensitive to proper tuning of the parameters that govern the data transport. Streams that are generated with "Autoset Internals" usually provide the optimal balance between performance and FPGA resource utilization.

In some rare cases, it may be required to make adjustments to the PCIe block's parameters in order to achieve the target bandwidth on host-to-FPGA streams. There is however no room for relevant changes in most FPGA families, as they are already set for best possible performance. Among Xilinx FPGAs, only Kintex-7 and Virtex-7 with a Gen2 PCIe block have any room for potential improvement.

In any case, this topic applies to Revision B, XL and XXL cores only, as revision A cores don't reach the data rates for which any improvement will be noticed.

The underlying issue is that data in the host-to-FPGA direction flows by virtue of DMA requests issued by the FPGA, which are fulfilled by the host by sending data. The delay between the requests and their respective data transmissions (completions) depends on the host's responsiveness to such requests, and varies from one processor architecture to another.

In order to make an effective use of the PCIe link's bandwidth when the maximal data rate is desired, several DMA requests are issued in parallel by the FPGA, thus ensuring that the host always has a request to handle. There is however a limitation on the number of requests in flight, which is imposed by the PCIe protocol's flow control. The limited resource, which is allocated by the flow control mechanism is called *completion credits*, and is configured for a PCIe endpoint. Generally speaking, more of these means more requests in flight are allowed, and also more FPGA resources used to implement the PCIe block.

The FPGA-to-host direction is much less affected, if at all, by credits allocation, as the FPGA sends the data along with the DMA requests. There is hence a little chance for an improvement by modifying the credit settings, as they have little influence.

The PCIe blocks in Xillybus bundles are configured for a full bandwidth utilization on mainstream x86-based desktop processors. Even though quite uncommon, it may be necessary to alter the configuration in order to attain the advertised host-to-FPGA bandwidth.

An improvement may be attained by increasing the number of completion credits (both header and data). This is done by invoking the configuration of the PCIe block IP in Vivado or Quartus, and modifying its parameters. For example, for a Kintex-7 configured in Vivado, this is done by selecting the "Core Capabilities" tab and setting the "BRAM Configuration Options". The "Perf level" is already set to the highest possible, so there's no room for improvement on this. However enabling "Buffering Optimized for Bus Mastering Application" increases the completion credits on the expense of the other credit types, which may the improve host-to-FPGA bandwidth without an adverse effect on the opposite direction.